

Course Project: Question Answering System

CS480: Database Systems

Fall 2025

1 Overview

In this project, you will design and implement a **document question-answering system** that combines two major parts:

1. **Relational Database Part:** You will design entities, relationships, and schemas for managing different entities of the system, like users, roles, documents, and logs. You will create an ER diagram, translate it into a relational schema (SQL), and implement CRUD operations on that.
2. **Vector Database Part:** You will implement a pipeline similar to **Retrieval-Augmented Generation (RAG)**.¹ Specifically, you will process unstructured documents by chunking them, generating embeddings, building vector indices, and running top- k retrieval queries. A language model will then generate answers with citations drawn from the retrieved passages.

1.1 Project Goals

- Learn ER diagram design and schema creation.
- Implement CRUD over a relational DB.
- Gain hands-on experience with document embedding vectors.
- Using popular vector indices (like IVF or HNSW).
- Generate answers with citations from retrieved passages.

Questions?

If you have any questions, feel free to contact the TA (Mohsen: mdehgh2@uic.edu) or visit during office hours on Wednesdays from 3:30–5:30 pm in CDRLC 2402.

2 Timeline & Phases

The project will be completed in several phases. Each phase has a deliverable that must be submitted on the due date. Late submissions will follow the course late policy.

¹What is RAG? See this blog post.

Phase 0: Teams (Deadline: Monday, September 15)

- Select a dataset (containing approximately 10-30 documents.)
- Submit the following:
 1. Description of the dataset and its source.
 2. List of team members (at most 3).
 3. Chosen team name.
 4. GitHub usernames of all members.

Phase 1: ER Model (Deadline: October 5th)

- Create an ER diagram for the system.
- The diagram should clearly include all required entities and their relationships.
- Submit screenshots or image files of the ER diagram.

Phase 2: Relational Schema (Deadline: October 22nd)

- Make sure to follow the requirements specified in Section ?? for your ER model.
 - You can still change/finalize your ER model in this phase.
- Translate the ER model (submitted in the previous phase of the project) into a relational schema, implemented as an SQL script using PostgreSQL's dialect.
- Submit just a `.sql` file that generates the relational schema.

Phase 3: Vector Pipeline (Deadline: November 17th)

1. Chunk the textual data from your dataset ($text \rightarrow chunks$).
2. Embed the chunks into vector representations ($chunks \rightarrow vectors$).
3. Store the vectors into a vector database ($vectors \rightarrow vectorDB$)
4. Answer a query by retrieving the nearest-neighbors in the stored vectors ($query$).
5. Commit all the code to the team's GitHub repository.
 - The *last commit* before November 17th (11:59 p.m.) is considered for grading this phase.
 - Look at *Section 4* for the details of deliverables.

Phase 4: Full System Integration (Deadline: November 30)

- See Section 3.

3 Full System Integration (100 pts)

This is the final phase of the project. In this phase, you are expected to complete the entire system according to the full set of requirements provided below.

Each group will deliver a 15-minute **demo presentation**. During the demo, you must walk through all required scenarios, and your grade for this phase will be based on the correctness and completeness of those demonstrations. More details about the demo will be provided in future.

Please ensure that all work is committed and pushed to the GitHub repository before the final deadline. **All group members** must make a meaningful contribution to the final codebase, and all members must participate in the demo presentation.

3.1 User Scenarios (30 pts)

1. A new user signs up and is successfully added to the system.
2. The newly registered user logs in using their credentials.
3. An admin logs in and retrieves a list of all registered users.
4. The admin edits the profile information of an existing user.
5. All user accounts are stored and retrieved from the SQL database.

3.2 Curator Scenarios (40 pts)

1. A curator uploads a new document to the system, and it becomes available for search.
2. A curator removes one of their previously uploaded documents from the system.
3. The system stores all document metadata correctly from the SQL database.
4. All documents and embeddings persist across system restarts and are automatically reloaded into the vector database.

3.3 EndUser Scenarios (30 pts)

1. An end user submits a query to the system and receives a response.
2. A curator uploads a new document, and the end user successfully queries information related to this newly added document.
3. All query logs are properly stored and retrievable from the SQL database.
4. An admin deletes an end user, and all associated data (e.g., query logs) is correctly removed from the system.

3.4 Bonus (up to 20 pts)

You can earn bonus points by creating a clean, intuitive, and visually appealing user interface for your system. This may include components such as login pages, document displays, or query-answering views—feel free to use your creativity in designing any UI elements you think enhance the user experience. Bonus points will be awarded competitively based on overall quality across groups. **At the end of the semester, a list of top-performing project groups will be highlighted on the course organization, and all repositories will be made public.**

4 Constraints & Tips

- Keep dataset modest to run locally.
- Use only public or course-provided content.
- Ensure reproducibility of your code (in GitHub Repository).
- If using an API, ensure there's a free tier or fallback template-only generator.

5 Suggested Tools

- **Relational Database:** PostgreSQL
- **ER Diagram Design:** draw.io, Lucidchart, dbdiagram.io
- **Document Parsing:** pdfminer.six, PyPDF2, textract, readability-lxml
- **Embeddings:** SentenceTransformers (e.g., `all-MiniLM-L6-v2`, `bge-small-en`), Hugging Face API models
- **Vector Index / Vector Database:** FAISS, pgvector (Postgres extension), Qdrant, Weaviate, Milvus
- **Backend / Application:** Python (CLI with argparse or click; minimal API with Flask or FastAPI)
- **LLM for Generation:**
 - OpenAI API (`gpt-4o-mini`, `gpt-4o`, or similar free-tier model)
 - Hugging Face Inference API (e.g., `mistralai/Mistral-7B-Instruct-v0.2`, `tiuae/falcon-7b-instruct`)
 - Local open-source models via `llama.cpp` or `Ollama` (if no API is used)